

Chapter 2

**Memory Management:  
Early Systems**

***Understanding Operating Systems,  
Fourth Edition***

# Objectives

You will be able to describe:

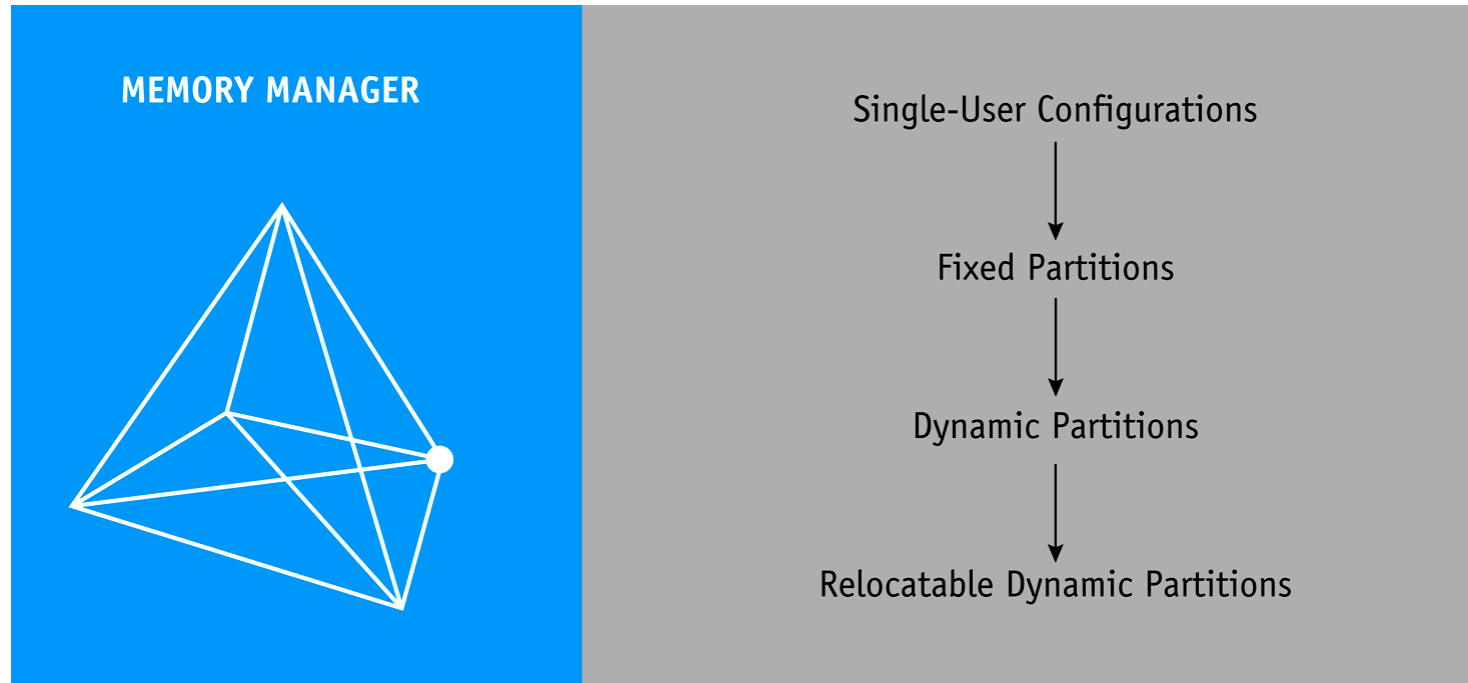
- The basic functionality of the three memory allocation schemes presented in this chapter: fixed partitions, dynamic partitions, relocatable dynamic partitions
- Best-fit memory allocation as well as first-fit memory allocation schemes
- How a memory list keeps track of available memory
- The importance of deallocation of memory in a dynamic partition system

# Objectives (continued)

Students should be able to describe:

- The importance of the bounds register in memory allocation schemes
- The role of compaction and how it improves memory allocation efficiency

# Memory Management: Early Systems



**“Memory is the primary and fundamental power, without which there could be no other intellectual operation.” —Samuel Johnson (1709–1784)**

# Memory Management: Early Systems

- **Types of memory allocation schemes:**
  - Single-user systems
  - Fixed partitions
  - Dynamic partitions
  - Relocatable dynamic partitions

# Single-User Contiguous Scheme

- **Single-User Contiguous Scheme:** Program is loaded in its entirety into memory and allocated as much contiguous space in memory as it needs
  - Jobs processed sequentially in single-user systems
  - Requires minimal work by the Memory Manager
    - Register to store the base address
    - Accumulator to keep track of the program size

# Single-User Contiguous Scheme (continued)

- **Disadvantages of Single-User Contiguous Scheme:**
  - Doesn't support multiprogramming
  - Not cost effective

# Fixed Partitions

- **Fixed Partitions:** Main memory is partitioned; one partition/job
  - Allows multiprogramming
  - Partition sizes remain static unless and until computer system is shut down, reconfigured, and restarted
  - Requires protection of the job's memory space
  - Requires matching job size with partition size




# Fixed Partitions (continued)

To allocate memory spaces to jobs, the operating system's Memory Manager must keep a table as shown below:

Partition Size	Memory Address	Access	Partition Status
100K	200K	Job 1	Busy
25K	300K	Job 4	Busy
25K	325K		Free
50K	350K	Job 2	Busy

Table 2.1: A simplified fixed partition memory table with the free partition shaded

# Fixed Partitions (continued)



**NOTE:** Job 3 must wait even though 70K of free space is available in Partition 1 where Job 1 occupies only 30K of the 100K available

Figure 2.1: Main memory use during fixed partition allocation of Table 2.1

# Fixed Partitions (continued)

- **Disadvantages:**
  - Requires entire program to be stored contiguously
  - Jobs are allocated space on the basis of first available partition of required size
  - Works well only if all of the jobs are of the same size or if the sizes are known ahead of time
  - Arbitrary partition sizes lead to undesired results
    - Too small a partition size results in large jobs having longer turnaround time
    - Too large a partition size results in memory waste or internal fragmentation

# Dynamic Partitions

- **Dynamic Partitions:** Jobs are given only as much memory as they request when they are loaded
  - Available memory is kept in contiguous blocks
  - Memory waste is comparatively small
- **Disadvantages:**
  - Fully utilizes memory only when the first jobs are loaded
  - Subsequent allocation leads to memory waste or external fragmentation

# Dynamic Partitions (continued)

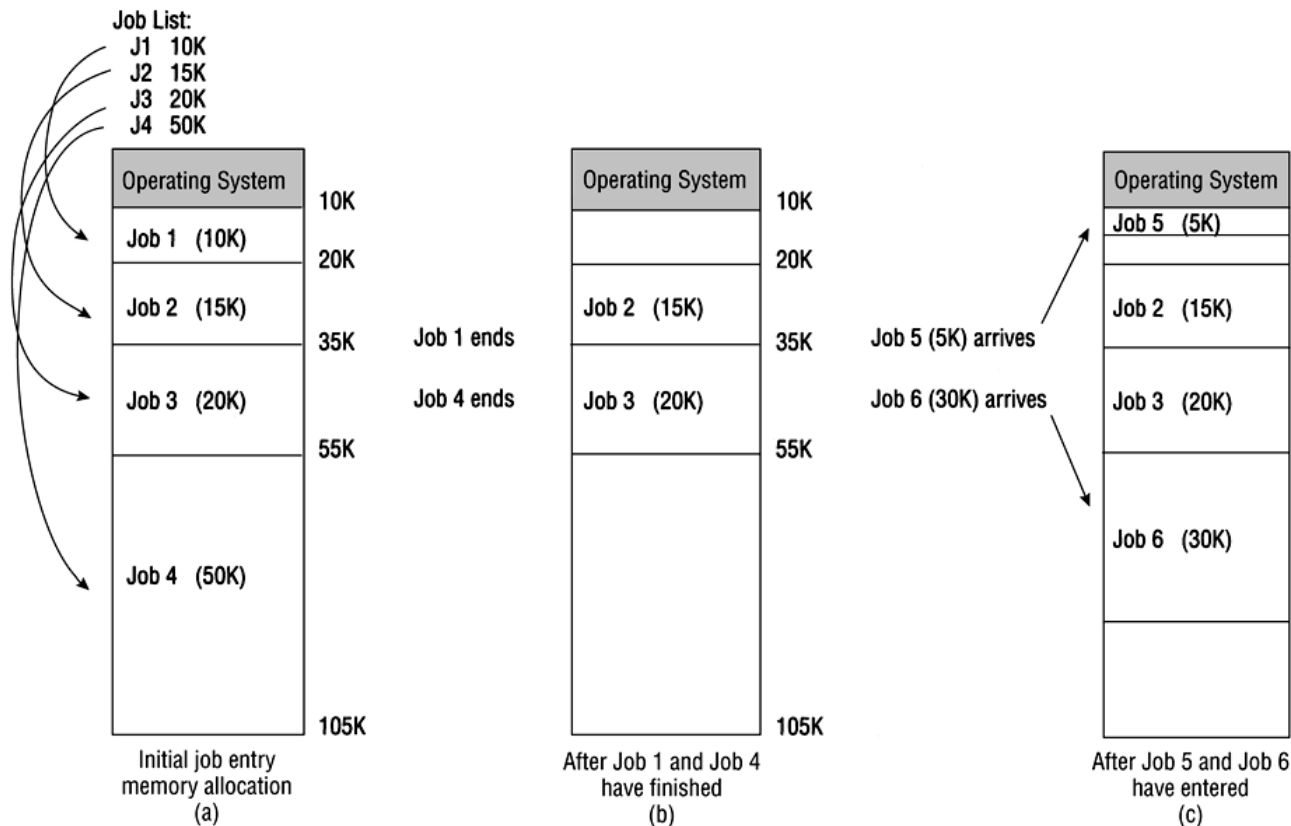


Figure 2.2: Main memory use during dynamic partition allocation

# Dynamic Partitions (continued)

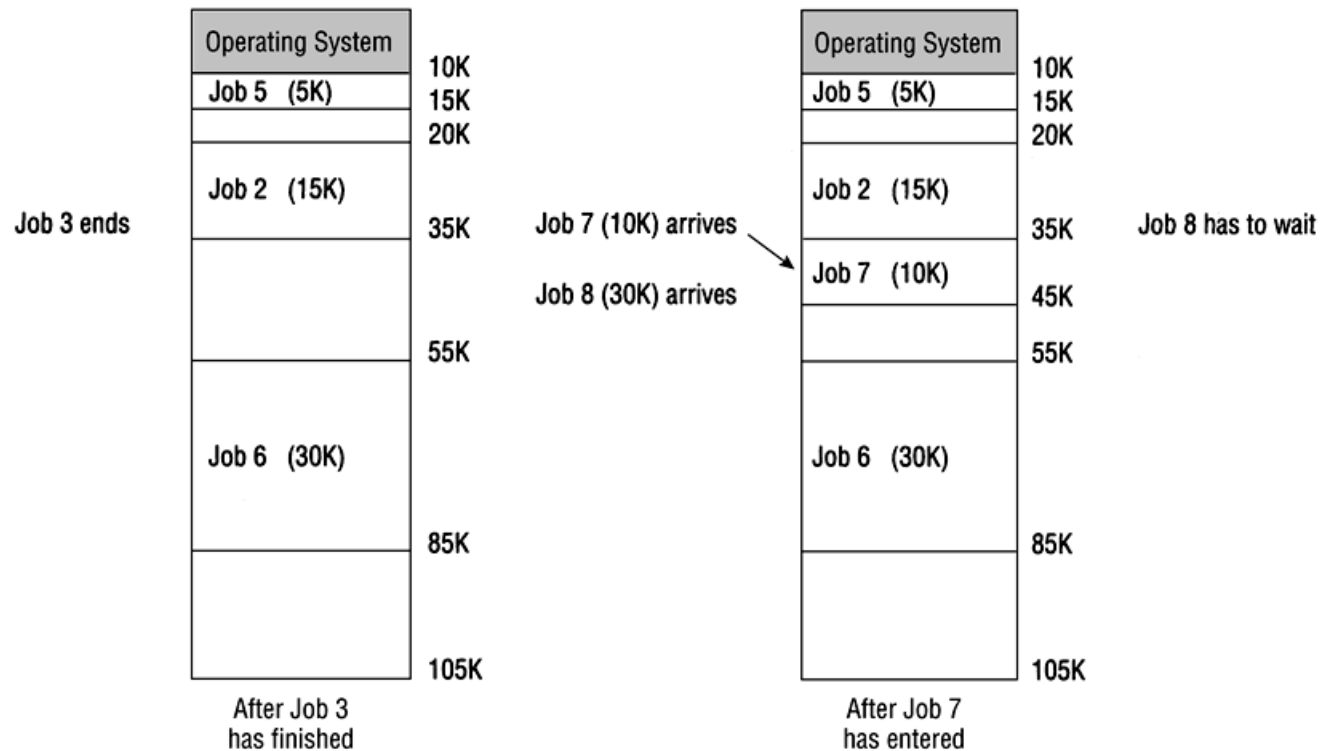


Figure 2.2 (continued): Main memory use during dynamic partition allocation

# Best-Fit Versus First-Fit Allocation

- Free partitions are allocated on the following basis:
  - **First-fit memory allocation:** First partition fitting the requirements
    - Leads to fast allocation of memory space
  - **Best-fit memory allocation:** Smallest partition fitting the requirements
    - Results in least wasted space
    - Internal fragmentation reduced but not eliminated

# Best-Fit Versus First-Fit Allocation (continued)

- **First-fit memory allocation:**
  - **Advantage:** Faster in making allocation
  - **Disadvantage:** Leads to memory waste
- **Best-fit memory allocation**
  - **Advantage:** Makes the best use of memory space
  - **Disadvantage:** Slower in making allocation



# Best-Fit Versus First-Fit Allocation (continued)

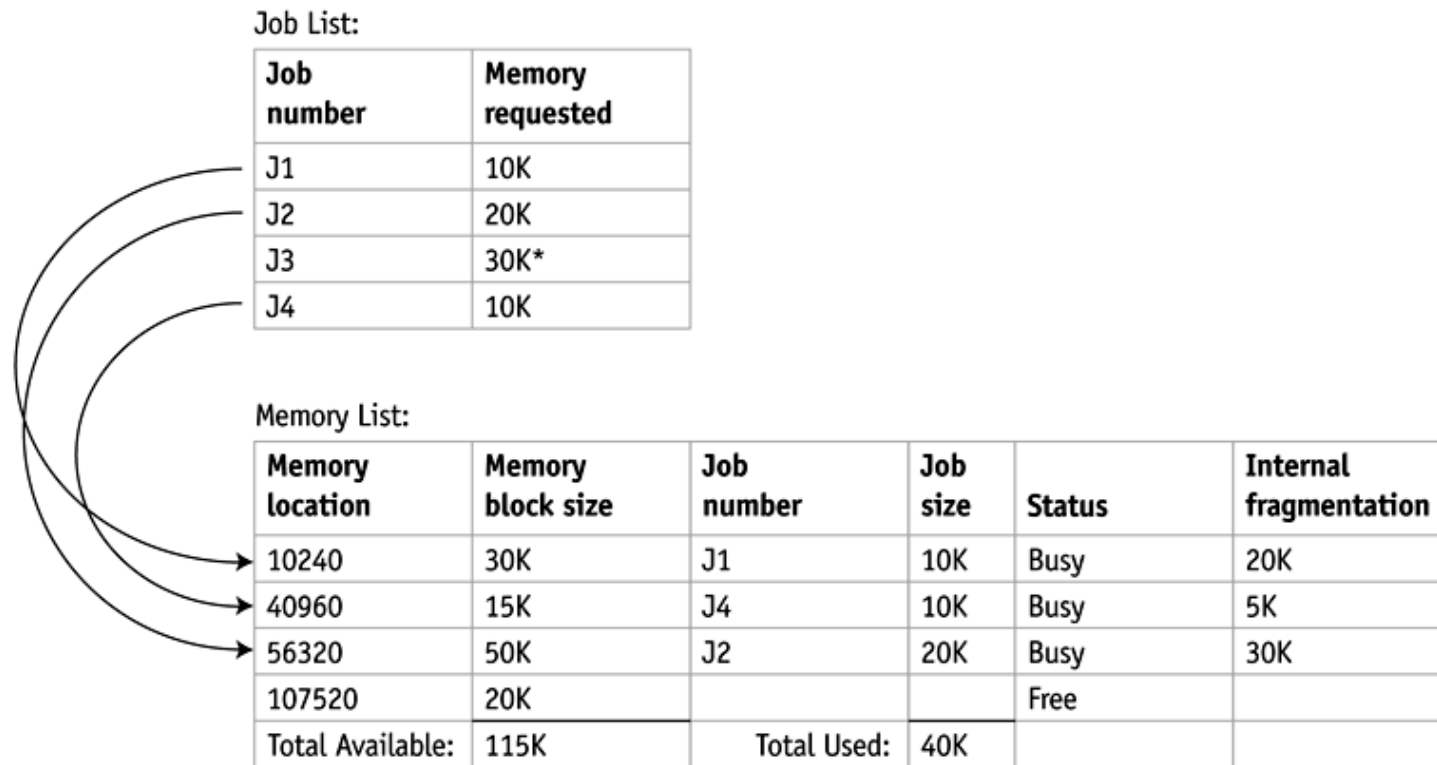


Figure 2.3: An example of a first-fit free scheme

# Best-Fit Versus First-Fit Allocation (continued)

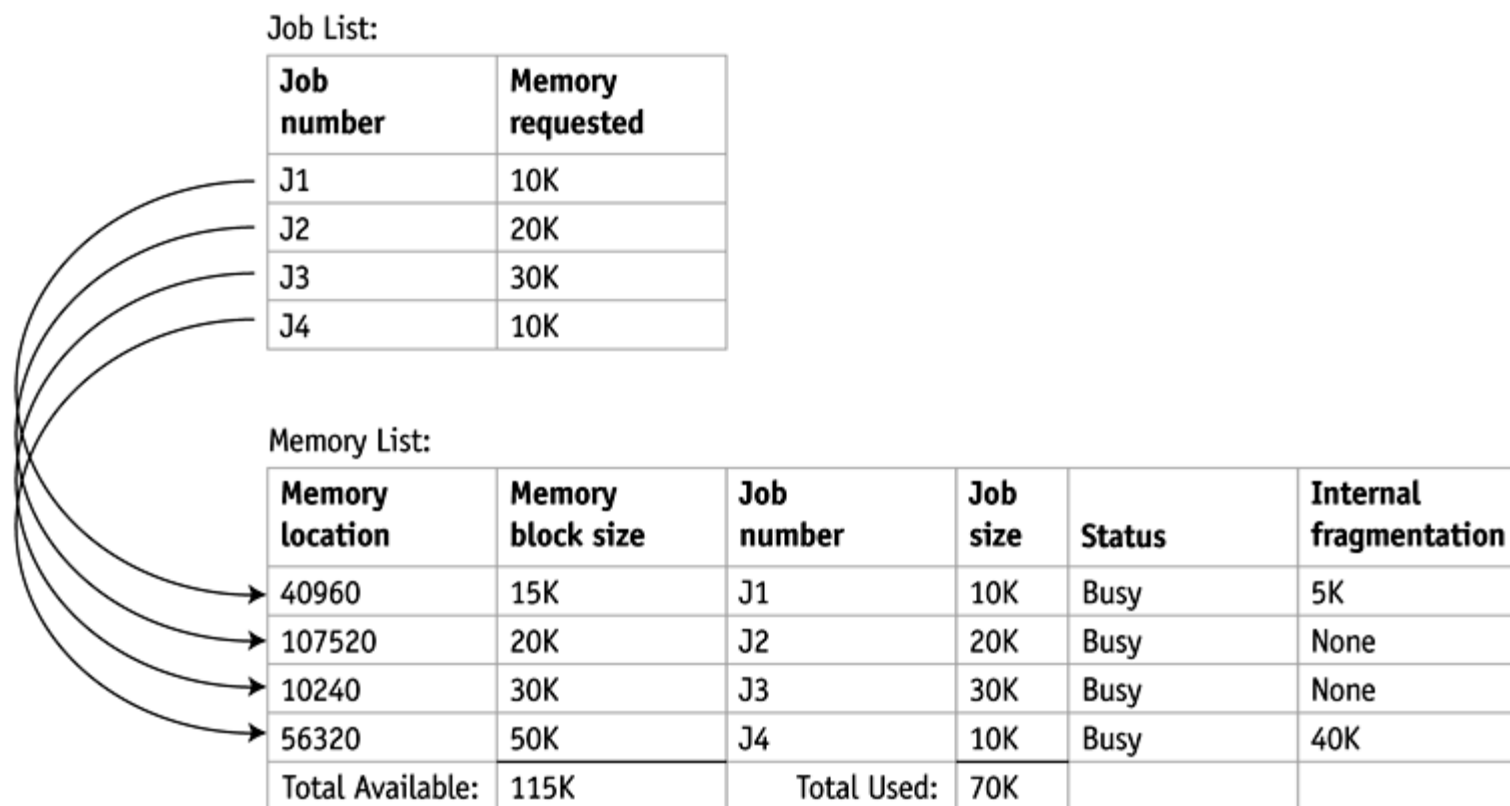


Figure 2.4: An example of a best-fit free scheme

# Best-Fit Versus First-Fit Allocation (continued)

- **Algorithm for First-Fit:**
  - Assumes Memory Manager keeps two lists, one for free memory and one for busy memory blocks
  - Loop compares the size of each job to the size of each memory block until a block is found that's large enough to fit the job
  - Job is stored into that block of memory
  - Memory Manager moves out of the loop to fetch the next job from the entry queue

# Best-Fit Versus First-Fit Allocation (continued)

- **Algorithm for First-Fit (continued):**
  - If the entire list is searched in vain, then the job is placed into a waiting queue
  - The Memory Manager then fetches the next job and repeats the process

# Best-Fit Versus First-Fit Allocation (continued)

Before Request		After Request	
Beginning Address	Memory Block Size	Beginning Address	Memory Block Size
4075	105	4075	105
5225	5	5225	5
6785	600	*6985	400
7560	20	7560	20
7600	205	7600	205
10250	4050	10250	4050
15125	230	15125	230
24500	1000	24500	1000

Table 2.2: Status of each memory block before and after a request is made for a block of 200 spaces using the first-fit algorithm

# Best-Fit Versus First-Fit Allocation (continued)

- **Algorithm for Best-Fit:**
  - Goal: find the smallest memory block into which the job will fit
  - Entire table must be searched before allocation

# Best-Fit Versus First-Fit Allocation (continued)

Before Request		After Request	
Beginning Address	Memory Block Size	Beginning Address	Memory Block Size
4075	105	4075	105
5225	5	5225	5
6785	600	6785	600
7560	20	7560	20
7600	205	*7800	5
10250	4050	10250	4050
15125	230	15125	230
24500	1000	24500	1000

Table 2.3: Status of each memory block before and after a request is made for a memory block of 200 spaces using the best-fit algorithm

# Best-Fit Versus First-Fit Allocation (continued)

- **Hypothetical allocation schemes:**
  - **Next-fit:** Starts searching from last allocated block, for the next available block when a new job arrives
  - **Worst-fit:** Allocates the largest free available block to the new job
    - Opposite of best-fit
    - Good way to explore the theory of memory allocation; might not be the best choice for an actual system



# Deallocation

- **Deallocation:** Freeing an allocated memory space
  - **For fixed-partition system:**
    - Straightforward process
    - When job completes, Memory Manager resets the status of the job's memory block to "free"
    - Any code—for example, binary values with 0 indicating free and 1 indicating busy—may be used

# Deallocation (continued)

- **For dynamic-partition system:**
  - Algorithm tries to combine free areas of memory whenever possible
  - Three cases:
    - **Case 1:** When the block to be deallocated is adjacent to another free block
    - **Case 2:** When the block to be deallocated is between two free blocks
    - **Case 3:** When the block to be deallocated is isolated from other free blocks

# Deallocation: Dynamic Partition System

- **Case 1: Joining Two Free Blocks**
  - Change list must reflect starting address of the new free block
    - In the example, 7600—which was the address of the first instruction of the job that just released this block
  - Memory block size for the new free space must be changed to show its new size—that is, the combined total of the two free partitions
    - In the example,  $(200 + 5)$

# Case 1: Joining Two Free Blocks

Beginning Address	Memory Block Size	Status
4075	105	Free
5225	5	Free
6785	600	Free
7560	20	Free
(7600)	(200)	(Busy) <sup>1</sup>
*7800	5	Free
10250	4050	Free
15125	230	Free
24500	1000	Free

<sup>1</sup>Although the numbers in parentheses don't appear in the free list, they've been inserted here for clarity. The job size is 200 and its beginning location is 7600.

Table 2.4: Original free list before deallocation for Case 1

# Case 1: Joining Two Free Blocks (continued)

Beginning Address	Memory Block Size	Status
4075	105	Free
5225	5	Free
6785	600	Free
7560	20	Free
*7600	205	Free
10250	4050	Free
15125	230	Free
24500	1000	Free

Table 2.5: Free list after deallocation for Case 1

# Deallocation:

## Dynamic Partition System (continued)

- **Case 2: Joining Three Free Blocks.** Deallocated memory space is between two free memory blocks
  - Change list to reflect the starting address of the new free block
    - In the example, 7560— which was the smallest beginning address
  - Sizes of the three free partitions must be combined
    - In the example,  $(20 + 20 + 205)$
  - Combined entry is given the status of null entry
    - In the example, 7600

## Case 2: Joining Three Free Blocks

Beginning Address	Memory Block Size	Status
4075	105	Free
5225	5	Free
6785	600	Free
*7560	20	Free
(7580)	(20)	(Busy) <sup>1</sup>
*7600	205	Free
10250	4050	Free
15125	230	Free
24500	1000	Free

<sup>1</sup> Although the numbers in parentheses don't appear in the free list, they have been inserted here for clarity.

Table 2.6: Original free list before deallocation for Case 2

## Case 2: Joining Three Free Blocks (continued)

Beginning Address	Memory Block Size	Status
4075	105	Free
5225	5	Free
6785	600	Free
7560	245	Free
*		(null entry)
10250	4050	Free
15125	230	Free
24500	1000	Free

Table 2.7: Free list after job has released memory



# Deallocation:

## Dynamic Partition System (continued)

- **Case 3: Deallocating an Isolated Block.** Space to be deallocated is isolated from other free areas
  - System learns that the memory block to be released is not adjacent to any free blocks of memory, it is between two other busy areas
  - Must search the table for a null entry
  - Null entry in the busy list occurs when a memory block between two other busy memory blocks is returned to the free list

## Case 3: Deallocating an Isolated Block

Beginning Address	Memory Block Size	Status
4075	105	Free
5225	5	Free
6785	600	Free
7560	245	Free
		(null entry)
10250	4050	Free
15125	230	Free
24500	1000	Free

Table 2.8: Original free list before deallocation for Case 3

## Case 3: Deallocating an Isolated Block (continued)

The job to be deallocated is of size 445 and begins at location 8805. The asterisk indicates the soon-to-be-free memory block.

Beginning Address	Memory Block Size	Status
7805	1000	Busy
*8805	445	Busy
9250	1000	Busy

Table 2.9: Memory list before deallocation

## Case 3: Deallocating an Isolated Block (continued)

Beginning Address	Memory Block Size	Status
7805	1000	Busy
*		(null entry)
9250	1000	Busy

Table 2.10: Busy list after the job has released its memory.  
The asterisk indicates the new null entry in the busy list.

## Case 3: Deallocating an Isolated Block (continued)

Beginning Address	Memory Block Size	Status
4075	105	Free
5225	5	Free
6785	600	Free
7560	245	Free
*8805	445	Free
10250	4050	Free
15125	230	Free
24500	1000	Free

Table 2.11: Free list after the job has released its memory.  
The asterisk indicates the new free block entry  
replacing the null entry

# Relocatable Dynamic Partitions

- **Relocatable Dynamic Partitions:**
  - Memory Manager relocates programs to gather together all of the empty blocks
  - Compact the empty blocks to make one block of memory large enough to accommodate some or all of the jobs waiting to get in

# Relocatable Dynamic Partitions (continued)

- **Compaction:** Reclaiming fragmented sections of the memory space
  - Every program in memory must be relocated so they are contiguous
  - Operating system must distinguish between addresses and data values
    - Every address must be adjusted to account for the program's new location in memory
    - Data values must be left alone

# Relocatable Dynamic Partitions (continued)

```
A      EXP 132, 144, 125, 110      ;the data values
BEGIN: MOVEI      1,0              ;initialize register 1
      MOVEI      2,0              ;initialize register 2
LOOP:  ADD        2,A(1)           ;add (A + reg 1) to reg 2
      ADDI        1,1              ;add 1 to reg 1
      CAIG        1,4,-1          ;is reg 1.4-1?
      JUMPA       LOOP            ;if not, go to Loop
      MOVE        3,2              ;if so, move reg 2 to reg 3
      IDIVI       3,4              ;divide reg 3 by 4,
                                   ;remainder to register 4
      EXIT                               ;end
      END
```

Figure 2.5: An assembly language program that performs a simple incremental operation



# Relocatable Dynamic Partitions (continued)

```
000000' 000000 000132      A:      EXP132,144,125,110
000001' 000000 000144
000002' 000000 000125
000003' 000000 000110

000004' 201 01 0 00 000000  BEGIN:  MOVEI      1,0
000005' 201 02 0 00 000000      MOVEI      2,0
000006' 270 02 0 01 000000'  LOOP:  ADD       2,A(1)
000007' 271 01 0 00 000001      ADDI      1,1
000008' 307 01 0 00 000003      CAIG      1,4-1
000009' 324 00 0 00 000006'      JUMPA     LOOP
000010' 200 03 0 00 000002      MOVE      3,2
000011' 231 03 0 00 000004      IDIVI     3,4
000012' 047 00 0 00 000012      EXIT

                                000000      END
```

Figure 2.6: The original assembly language program  
after it has been processed by the  
assembler

# Relocatable Dynamic Partitions (continued)

- Compaction issues:
  - What goes on behind the scenes when relocation and compaction take place?
  - What keeps track of how far each job has moved from its original storage area?
  - What lists have to be updated?

# Relocatable Dynamic Partitions (continued)

- **What lists have to be updated?**
  - Free list must show the partition for the new block of free memory
  - Busy list must show the new locations for all of the jobs already in process that were relocated
  - Each job will have a new address except for those that were already at the lowest memory locations

# Relocatable Dynamic Partitions (continued)

- Special-purpose registers are used for relocation:
  - **Bounds register**
    - Stores highest location accessible by each program
  - **Relocation register**
    - Contains the value that must be added to each address referenced in the program so it will be able to access the correct memory addresses after relocation
    - If the program isn't relocated, the value stored in the program's relocation register is zero

# Relocatable Dynamic Partitions (continued)

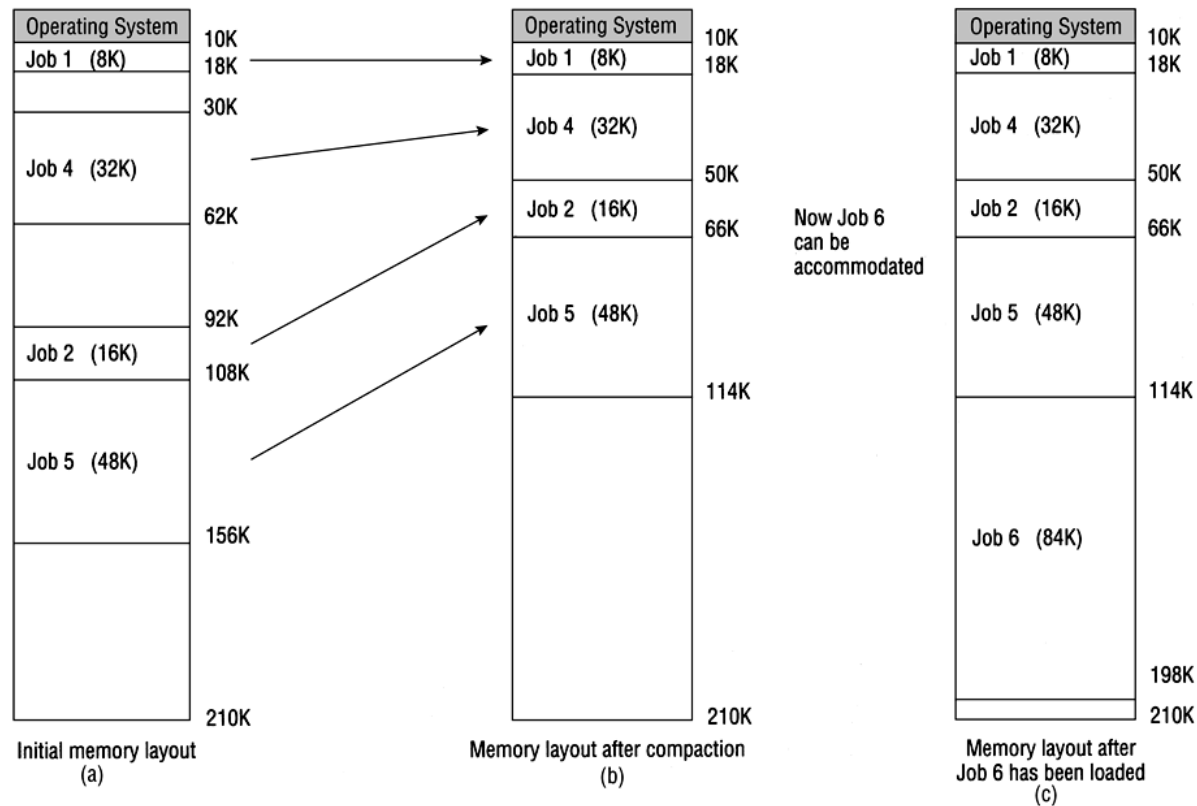


Figure 2.7: Three snapshots of memory before and after compaction

# Relocatable Dynamic Partitions (continued)

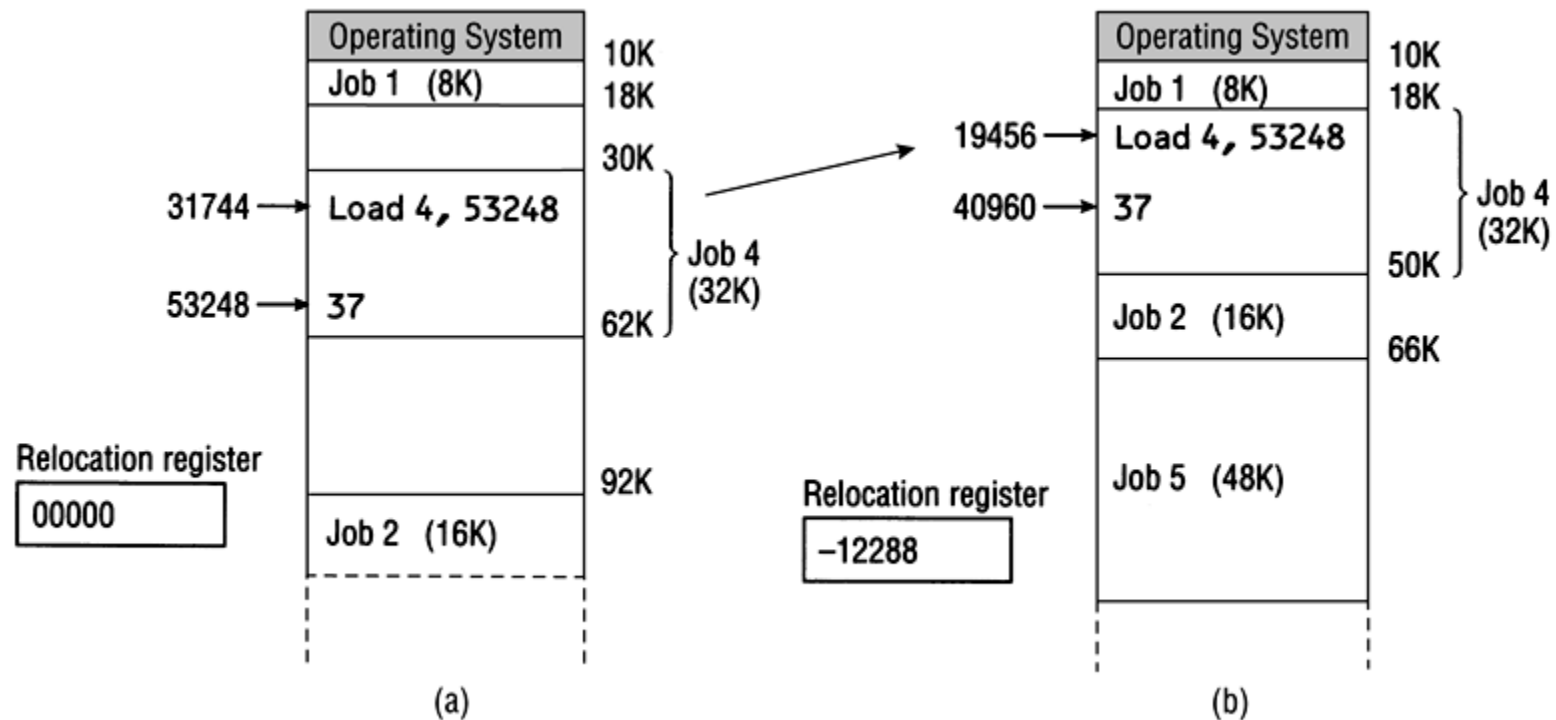


Figure 2.8: Contents of relocation register and close-up of Job 4 memory area (a) before relocation and (b) after relocation and compaction

# Relocatable Dynamic Partitions (continued)

- Compacting and relocating optimizes the use of memory and thus improves throughput
  - Options for when and how often it should be done:
    - When a certain percentage of memory is busy
    - When there are jobs waiting to get in
    - After a prescribed amount of time has elapsed
- Goal:** Optimize processing time and memory use while keeping overhead as low as possible

# Summary

- Four memory management techniques were used in early systems: single-user systems, fixed partitions, dynamic partitions, and relocatable dynamic partitions
- Memory waste in dynamic partitions is comparatively small as compared to fixed partitions
- First-fit is faster in making allocation but leads to memory waste
- Best-fit makes the best use of memory space but slower in making allocation



## Summary (continued)

- Compacting and relocating optimizes the use of memory and thus improves throughput
- All techniques require that the entire program must:
  - Be loaded into memory
  - Be stored contiguously
  - Remain in memory until the job is completed
- Each technique puts severe restrictions on the size of the jobs: can only be as large as the largest partitions in memory