

# Chapter 4

## **Processor Management**

***Understanding Operating Systems,  
Fourth Edition***

# Objectives

You will be able to describe:

- The difference between job scheduling and process scheduling, and how they relate
- The advantages and disadvantages of process scheduling algorithms that are preemptive versus those that are nonpreemptive
- The goals of process scheduling policies
- Up to six different process scheduling algorithms
- The role of internal interrupts and the tasks performed by the interrupt handler

# Overview

- **Program (Job):**
  - A unit of work that has been submitted by user to an operating system
  - An inactive unit, such as a file stored on a disk
- **Process (Task):**
  - An active entity, which requires a set of resources, including a processor and special registers, to perform its function
  - A single instance of an executable program

# Overview (continued)

- **Processor (CPU):** performs calculations and executes programs
  - **In single-user systems:**
    - Processor is busy only when user is executing a job, at all other times it is idle
    - Processor management is simple
  - **In a multiprogramming environment:**
    - Processor must be allocated to each job in a fair and efficient manner
    - Requires scheduling policy and a scheduling algorithm

# Overview (continued)

- **Interrupt:** A hardware signal that suspends execution of a program and activates the execution of interrupt handler
- **Context Switch:** Saving a job's processing information in its PCB when interrupted
  - Context switching occurs in all preemptive policies

# Job Scheduling Versus Process Scheduling

- Processor Manager has two submanagers:
  - **Job Scheduler:**
    - In charge of job scheduling
    - Initiates the job based on certain criteria
  - **Process Scheduler:**
    - In charge of process scheduling
    - Assigns the CPU to execute processes of those jobs placed on READY queue by Job Scheduler

# Job Scheduling Versus Process Scheduling (continued)

## **Job Scheduler** (High level scheduler):

- Initiates the job based on certain criteria
- Puts jobs in a sequence that uses all system's resources as fully as possible
- Strives for balanced mix of jobs with large I/O interaction and jobs with lots of computation
- Tries to keep most system components busy most of time

# Job Scheduling Versus Process Scheduling (continued)

## **Process Scheduler** (Low level scheduler):

- Determines which jobs will get the CPU, when, and for how long
- Decides when processing should be interrupted
- Determines which queues the job should be moved to during its execution
- Recognizes when a job has concluded and should be terminated



# Job Scheduling Versus Process Scheduling (continued)

- **I/O-bound jobs** have many brief CPU cycles and long I/O cycles, e.g., printing a series of documents
- **CPU-bound jobs** have long CPU cycles and shorter I/O cycles, e.g., finding the first 300 prime numbers
- Total effect of all CPU cycles, from both I/O-bound and CPU-bound jobs, approximates a **Poisson distribution curve**

# Job Scheduling Versus Process Scheduling (continued)




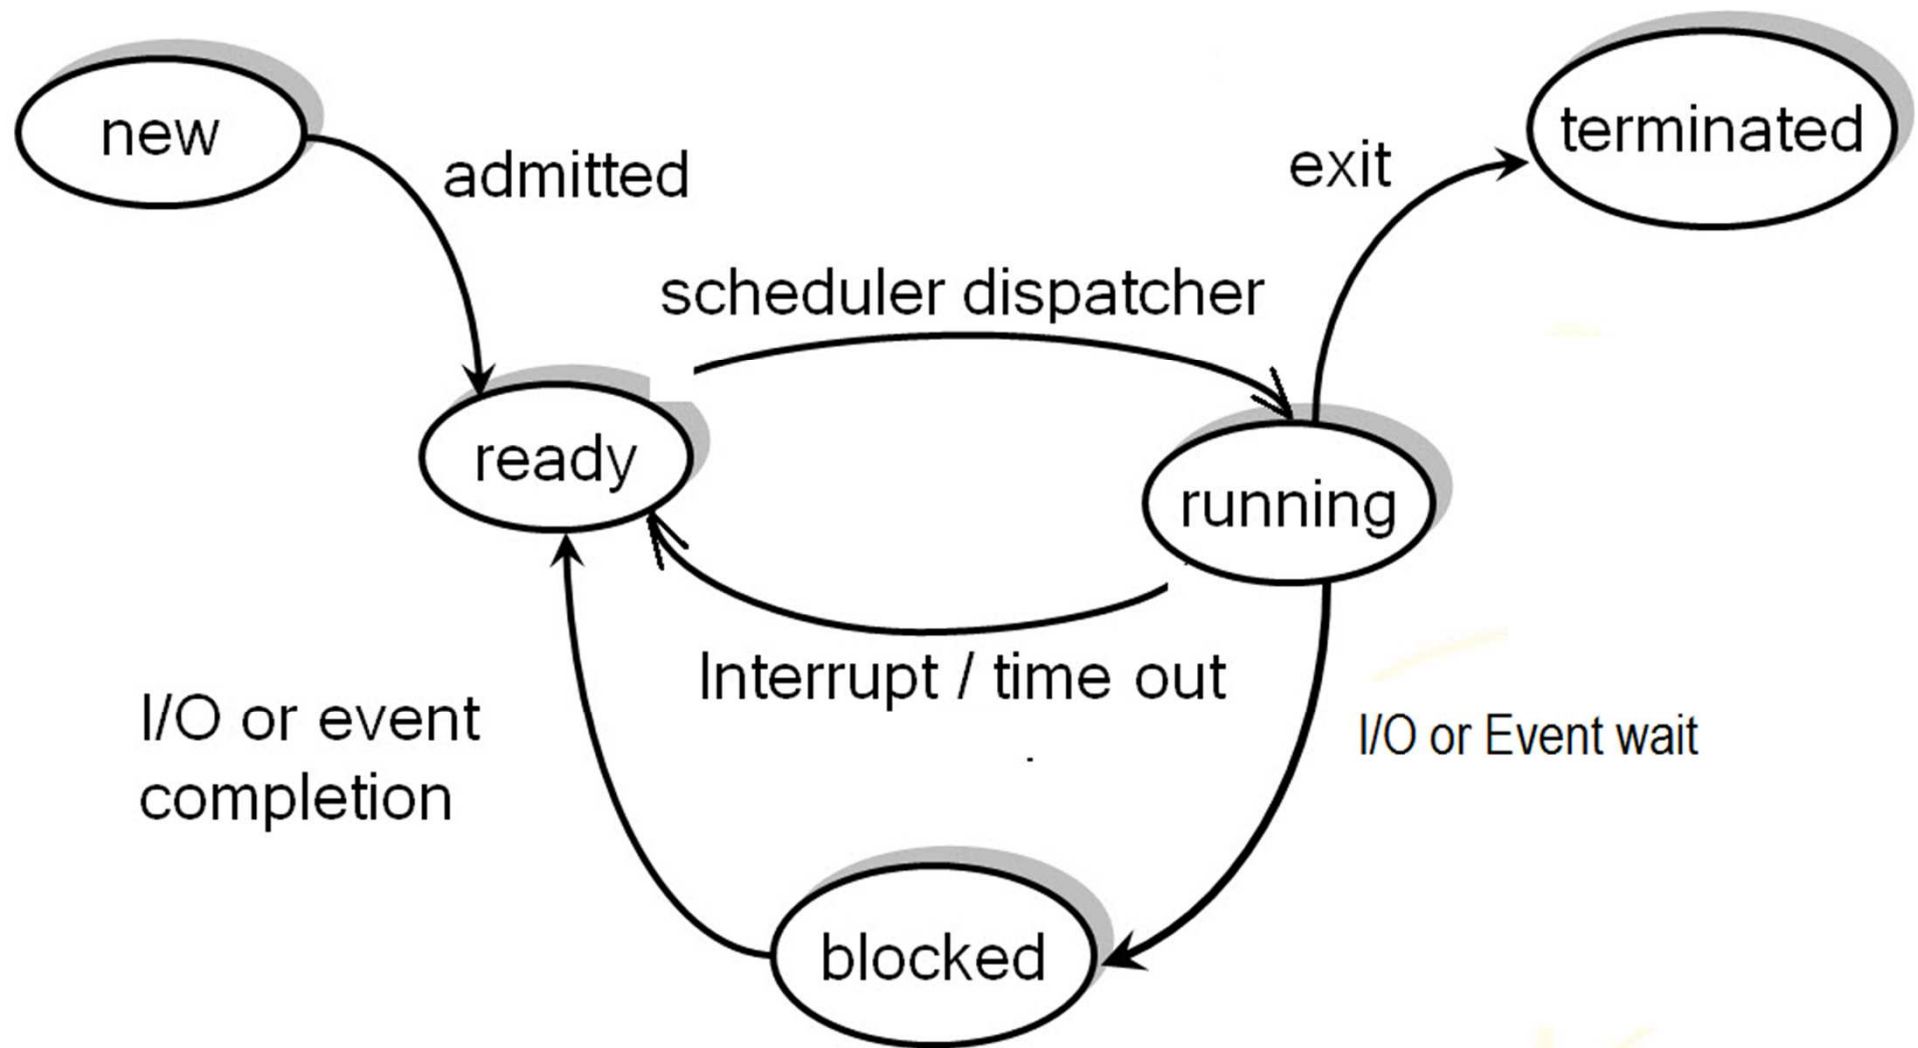
Figure 4.1: Distribution of CPU cycle times

# Job Scheduling Versus Process Scheduling (continued)

- **Middle level scheduler** (third layer): Used in a highly interactive environment
  - Removes active jobs from memory to reduce degree of multiprogramming
  - Allows jobs to be completed faster

# Job and Process Status

- **Job status:** A job takes one of the following states as it moves through the system
  - HOLD
  - READY
  - WAITING
  - RUNNING
  - FINISHED



# Job and Process Status (continued)

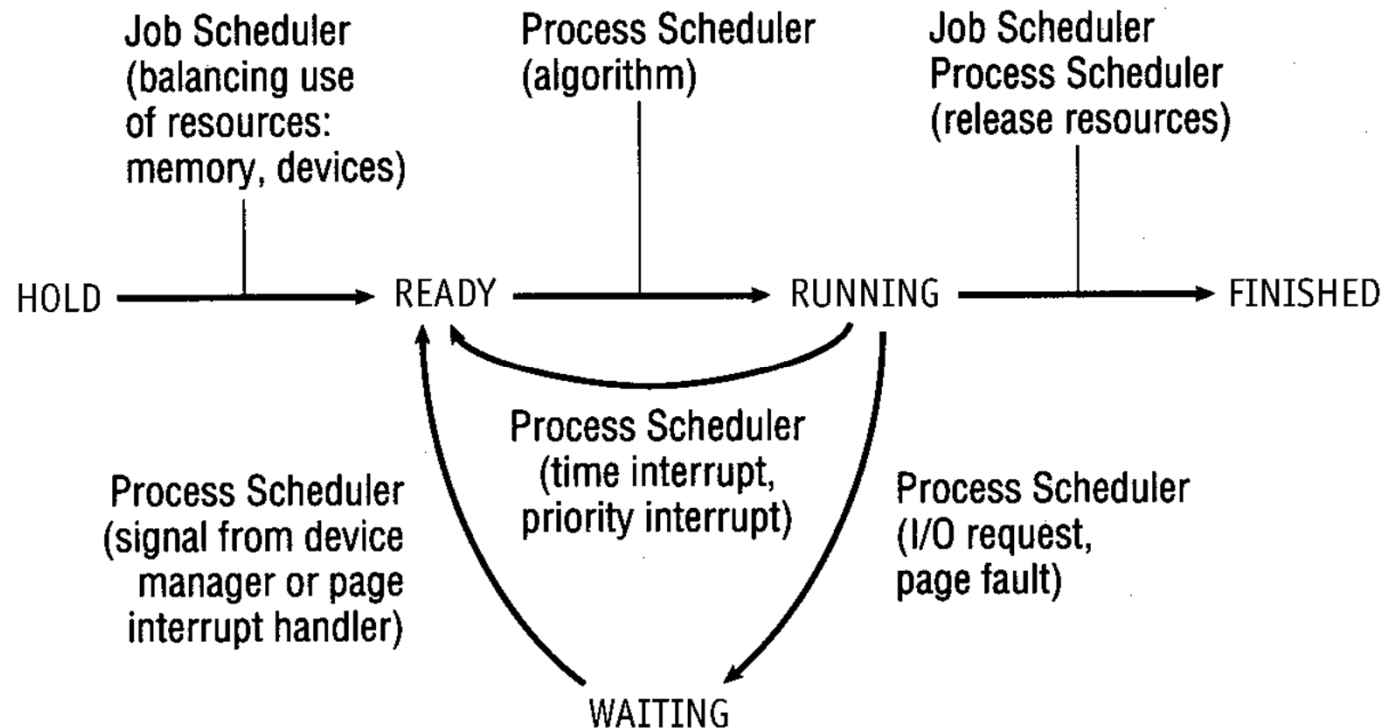


Figure 4.2: A typical job (or process) changes status as it moves through the system from HOLD to FINISHED

# Job and Process Status (continued)

- **Transition** from one status to another is initiated by either the Job Scheduler (**JS**) or the Process Scheduler (**PS**):
  - HOLD to READY: JS, using a predefined policy
  - READY to RUNNING: PS, using some predefined algorithm
  - RUNNING back to READY: PS, according to some predefined time limit or other criterion
  - RUNNING to WAITING: PS, and is initiated by an instruction in the job

# Job and Process Status (continued)

- **Transition** (continued):
  - WAITING to READY: PS, and is initiated by signal from I/O device manager that I/O request has been satisfied and job can continue.
  - RUNNING to FINISHED: PS or JS, if job is finished or error has occurred



# Process Control Blocks

- **Process Control Block (PCB):** Data structure that contains basic info about the job including
  - What it is
  - Where it's going
  - How much of its processing has been completed
  - Where it's stored
  - How much it has spent in using resources

# Process Control Blocks (continued)

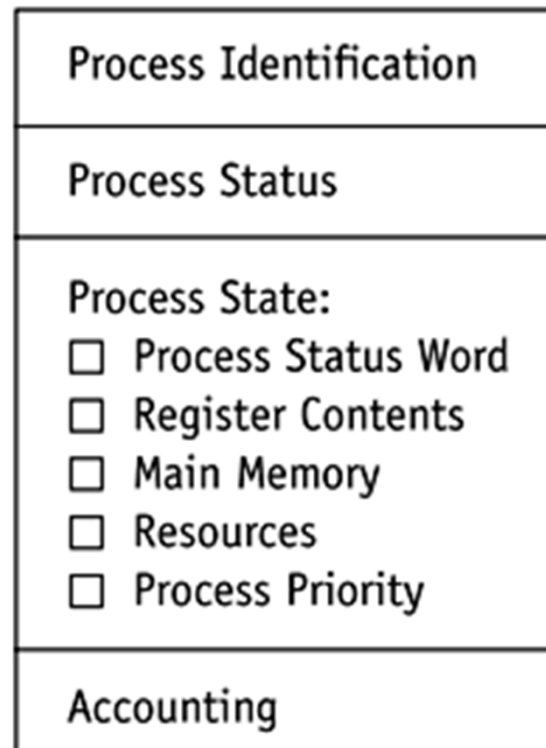


Figure 4.3: Contents of each job's Process Control Block

# Process Control Blocks (continued)

- **Contents of Process Control Block (PCB):**
  - **Process identification**
  - **Process status** (HOLD, READY, RUNNING, WAITING)
  - **Process state** (process status word, register contents, main memory info, resources, process priority)
  - **Accounting** (CPU time, total time, memory occupancy, I/O operations, number of input records read, etc.)

# PCBs and Queuing

- **PCB of a job:** Contains all of the data about the job needed by the operating system to manage the processing of the job
  - Created when job scheduler accepts the job
  - Updated as job goes from beginning to end of its execution
- **Queues** use PCBs to track jobs
  - PCBs, not jobs, are linked to form queues
  - Queues must be managed by process scheduling policies and algorithms

# PCBs and Queuing (continued)

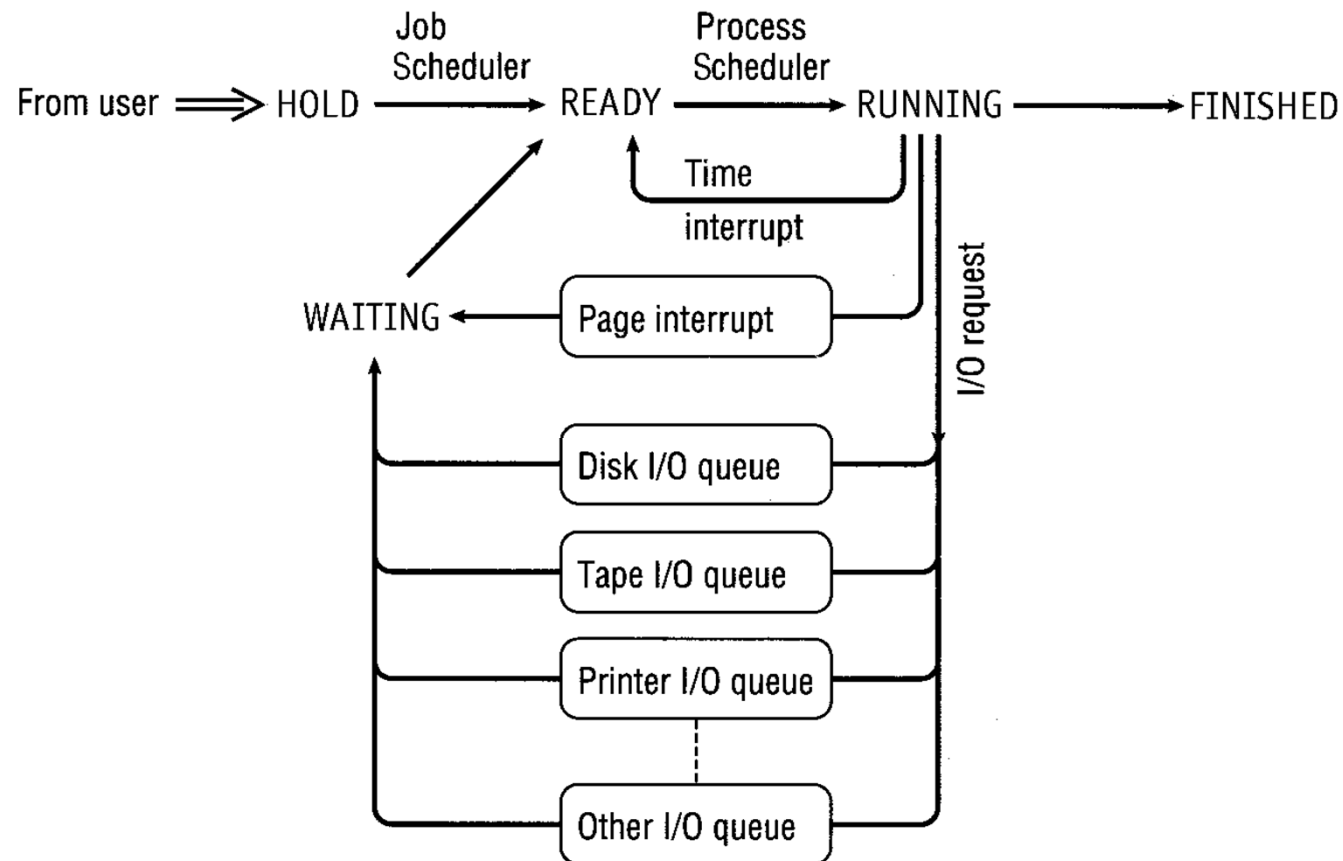


Figure 4.4: Queuing paths from HOLD to FINISHED

# Process Scheduling Policies

- Operating system must resolve three limitations of a system before scheduling all jobs in a multi-programming environment:
  - Finite number of resources (e.g., disk drives, printers, and tape drives)
  - Some resources can't be shared once they're allocated (e.g., printers)
  - Some resources require operator intervention (e.g., tape drives)

# Process Scheduling Policies (continued)

- A **good process scheduling policy** should
- **Maximize throughput** by running as many jobs as possible in a given amount of time
  - **Minimize response time** by quickly turning around interactive requests
  - **Minimize turnaround time** by moving entire jobs in and out of system quickly
  - **Minimize waiting time** by moving jobs out of READY queue as quickly as possible

# Process Scheduling Policies (continued)

- (continued):
  - **Maximize CPU efficiency** by keeping CPU busy 100 percent of time
  - **Ensure fairness** for all jobs by giving every one an equal amount of CPU and I/O time



# Process Scheduling Policies (continued)

- **Need for Interrupts:** When a job claims CPU for a very long time before issuing an I/O request
  - Builds up READY queue & empties I/O queues
  - Creates an unacceptable imbalance in the system
- Process Scheduler uses interrupts when a predetermined slice of time has expired
  - Suspends all activity on the currently running job
  - Reschedules it into the READY queue

# Process Scheduling Policies (continued)

- **Types of Scheduling Policies:**
  - **Preemptive scheduling policy:**
    - Interrupts processing of a job and transfers the CPU to another job
  - **Nonpreemptive scheduling policy:**
    - Functions without external interrupts

# Process Scheduling Algorithms

- **Types of Process Scheduling Algorithms:**
  - First Come, First Served (FCFS)
  - Shortest Job Next (SJN)
  - Priority Scheduling
  - Shortest Remaining Time (SRT)
  - Round Robin
  - Multiple Level Queues


# First-Come, First-Served

- Nonpreemptive
- Handles jobs according to their arrival time: the earlier they arrive, the sooner they're served
- Simple algorithm to implement: uses a FIFO queue
- Good for batch systems; unacceptable for interactive systems
- Turnaround time is unpredictable

# First-Come, First-Served (continued)

Jobs arrival sequence: A, B, C

- Job A has a CPU cycle of 15 milliseconds
- Job B has a CPU cycle of 2 milliseconds
- Job C has a CPU cycle of 1 millisecond




Average turnaround time: 16.67 s

Figure 4.5: Timeline for job sequence A, B, C using the FCFS algorithm

# First-Come, First-Served (continued)

Jobs arrival sequence: C, B, A

- Job A has a CPU cycle of 15 milliseconds
- Job B has a CPU cycle of 2 milliseconds
- Job C has a CPU cycle of 1 millisecond



Average turnaround time: 7.3 s

Figure 4.6: Timeline for job sequence C, B, A using the FCFS algorithm

# Shortest Job Next (SJN)


- Nonpreemptive
- Handles jobs based on length of their CPU cycle time
- Easiest to implement in batch environments
- Doesn't work in interactive systems
- Optimal only when all jobs are available at same time and the CPU estimates are available and accurate

# Shortest Job Next (continued)

Four batch jobs A, B, C, D, all in the READY queue

Job:            A   B   C   D

CPU cycle: 5   2   6   4



Average turnaround time: 9 s

Figure 4.7: Timeline for job sequence B, D, A, C using the SJN algorithm



# Priority Scheduling

- Nonpreemptive
- Gives preferential treatment to important jobs
  - Programs with highest priority are processed first
  - Not interrupted until CPU cycles are completed or a natural wait occurs
- FCFS policy is used if two or more jobs with equal priority in READY queue
- System administrator or Processor Manager use different methods of assigning priorities

# Shortest Remaining Time

- Preemptive version of the SJN algorithm
- Processor allocated to job closest to completion
  - Current job can be preempted if newer job in READY queue has shorter time to completion
- Cannot be implemented in interactive system
  - Requires advance knowledge of the CPU time required to finish each job
- SRT involves more overhead than SJN
  - OS monitors CPU time for all jobs in READY queue and performs context switching

# Shortest Remaining Time (continued)

Arrival time:	0	1	2	3
Job:	A	B	C	D
CPU cycle:	6	3	1	4

Job:	A	B	C	D
Turnaround:	14	4	1	6
Average Turnaround:	6.25s			




Figure 4.8: Timeline for job sequence A, B, C, D using the preemptive SRT algorithm

# Shortest Remaining Time (continued)

Arrival time:	0	1	2	3
Job:	A	B	C	D
CPU cycle:	6	3	1	4

Job:	A	B	C	D
Turnaround:	6	9	5	11
Average Turnaround:	7.75s			




Figure 4.9: Timeline for job sequence A, B, C, D using the nonpreemptive SJN algorithm

# Round Robin

- Preemptive
- Used extensively in interactive systems
- Based on a predetermined slice of time (time quantum) that's given to each job
- Size of time quantum crucial to system performance
  - Usually varies from 100 ms to 1-2 s
- Ensures CPU is equally shared among all active processes and is not monopolized by any one job

# Round Robin (continued)

Arrival time:	0	1	2	3
Job:	A	B	C	D
CPU cycle:	8	4	9	5

Job:	A	B	C	D
Turnaround:	20	7	24	22
Average Turnaround:	18.25 s			

Time slice: 4ms




Figure 4.10: Timeline for job sequence A, B, C, D using the preemptive round robin algorithm

## Round Robin (continued)

- **If Job's CPU cycle > time quantum**
  - Job is preempted and put at the end of the READY queue and its information is saved in its PCB
- **If Job's CPU cycle < time quantum**
  - If job is finished, all resources allocated to it are released & completed job is returned to user
  - If interrupted by I/O request, then info is saved in PCB & it is linked at end of the appropriate I/O queue
  - Once I/O request is satisfied, job returns to end of READY queue to await allocation of CPU

# Round Robin (continued)




Figure 4.11: Context switches for job A with three different time quanta. In (a) the job finishes before the time quantum expires. In (b) and (c), the time quantum expires first, interrupting the job.



## Round Robin (continued)

- **Efficiency** depends on the size of time quantum in relation to the average CPU cycle
- **If the quantum is too large** - larger than most CPU cycles
  - Algorithm reduces to the FCFS scheme
- **If the quantum is too small**
  - Amount of context switching slows down the execution of the jobs
  - Amount of overhead is dramatically increased

## Round Robin (continued)

- General rules of thumb for selecting the proper time quantum:
  - Should be long enough to allow 80% of CPU cycles to run to completion
  - Should be at least 100 times longer than the time required to perform one context switch
- These rules are flexible and depend on the system

# Multiple-Level Queues

- Work in conjunction with several other schemes
- Found in systems with jobs that can be grouped according to a common characteristic
- Examples:
  - Priority-based system with different queues for each priority level
  - System with all CPU-bound jobs in one queue and all I/O-bound jobs in another
  - Hybrid system with batch jobs in background queue and interactive jobs in a foreground queue

# Multiple-Level Queues (continued)

- Four primary methods to the movement of jobs:
  - No Movement Between Queues
  - Movement Between Queues
  - Variable Time Quantum Per Queue
  - Aging

# Multiple-Level Queues (continued)

- **No Movement Between Queues:**
  - The processor is allocated to the jobs in the high-priority queue in FCFS fashion
  - Allocated to jobs in lower priority queues only when the high priority queues are empty
- **Movement Between Queues**
  - Adjusts the priorities assigned to each job
  - A job may also have its priority increased
  - Good in interactive systems

# Multiple-Level Queues (continued)

- **Variable Time Quantum Per Queue:**
  - Each of the queues is given a time quantum twice as long as the previous queue
  - CPU-bound job can execute for longer and longer periods of time, thus improving its chances of finishing faster
- **Aging:**
  - System moves the old job to the next highest queue, and so on until it reaches the top queue
  - Ensures that jobs in the lower-level queues will eventually complete their execution

# A Word About Interrupts

- **Types of Interrupts:**
  - Page interrupts to accommodate job requests
  - Time quantum expiration interrupts
  - I/O interrupts when READ or WRITE command is issued
  - Internal interrupts (synchronous interrupts) result from arithmetic operation or job instruction
  - Illegal arithmetic operations (e.g., dividing by 0).
  - Illegal job instructions (e.g., attempts to access protected storage locations)

# A Word About Interrupts (continued)

- **Interrupt handler:** Control program that handles the interruption sequence of events
- When operating system detects a nonrecoverable error, the interrupt handler follows this sequence:
  - The type of interrupt is described and stored
  - The state of the interrupted process is saved
  - The interrupt is processed
  - The processor resumes normal operation



# Summary

- Process scheduler assigns the CPU to execute processes of those jobs placed on READY queue by Job Scheduler
- Total effect of all CPU cycles, from both I/O-bound and CPU-bound jobs, approximates a Poisson distribution curve
- Transition from one status to another is initiated by either Job Scheduler (JS) or Process Scheduler (PS)
- PCB of a job contains all data about the job needed by OS to manage the processing of the job

## Summary (continued)

- A good process scheduling policy should maximize CPU efficiency by keeping CPU busy 100 percent of time
- FIFO has simple algorithm to implement but turnaround time is unpredictable
- SJN minimizes average waiting time but results in infinite postponement of some jobs
- Priority scheduling ensures fast completion of important jobs but results in infinite postponement of some jobs

## Summary (continued)

- SRT ensures fast completion of short jobs but involves more overhead than SJN, incurred by context switching
- Efficiency in round robin policy depends on the size of time quantum in relation to the average CPU cycle
- Multiple-level queues counteract indefinite postponement with aging or other queue movement